

```
In [2]: import numpy as np

import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Activation, Dense, Input

import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```


```
In [4]: x_train.shape
```

```
Out[4]: (60000, 28, 28)
```

```
In [7]: print("Actual Lable :", y_train[80])
plt.imshow(x_train[80], cmap="gray")
```

Actual Lable : 9

```
Out[7]: <matplotlib.image.AxesImage at 0x2017e0f1240>
```



**Create model**

```
In [8]: def create_model_1():
        model = Sequential() # MLP Model
        model.add(Dense(20, activation="sigmoid", input_shape=(28*28,)))
        model.add(Dense(32, activation="sigmoid"))
        model.add(Dense(10, activation="softmax"))
        model.compile(loss='categorical_crossentropy', optimizer='SGD', metrics=["accuracy"])
        model.summary()
        return model

model_1 = create_model_1()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_3 (Dense)	(None, 20)	15700
dense_4 (Dense)	(None, 32)	672
dense_5 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 16,702		
Trainable params: 16,702		
Non-trainable params: 0		
=====	=====	=====

```
In [9]: y_train.shape
```

```
Out[9]: (60000,)
```

## Covert lables to one-hot vectors for MLP model training

```
In [11]: from tensorflow.keras.utils import to_categorical
```

```
In [12]: y_train_vecs = to_categorical(y_train, num_classes=10)
        y_test_vecs = to_categorical(y_test, num_classes=10 )
        print(y_train_vecs.shape, y_test_vecs.shape)
```

```
(60000, 10) (10000, 10)
```

```
In [13]: # Inspect vectors  
y_test_vecs[0], y_test[0]
```

```
Out[13]: (array([0., 0., 0., 0., 0., 0., 0., 1., 0., 0.], dtype=float32), 7)
```

**Let's train the model**

```
In [14]: history = model_1.fit(x=x_train.reshape((60000, 28*28)),
                                y=y_train_vecs,
                                batch_size=64,
                                epochs=20,
                                verbose=1,
                                validation_data=(x_test.reshape((10000, 28*28)), y_test_vecs))
```

Train on 60000 samples, validate on 10000 samples

WARNING:tensorflow:From C:\Users\subhendu\_mishra\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Epoch 1/20

60000/60000 [=====] - 5s 78us/sample - loss: 2.0916 - acc: 0.4632 - val\_loss: 1.8798  
- val\_acc: 0.6449

Epoch 2/20

60000/60000 [=====] - 5s 80us/sample - loss: 1.6939 - acc: 0.6765 - val\_loss: 1.5100  
- val\_acc: 0.7030

Epoch 3/20

60000/60000 [=====] - 4s 71us/sample - loss: 1.3540 - acc: 0.7287 - val\_loss: 1.2133  
- val\_acc: 0.7566

Epoch 4/20

60000/60000 [=====] - 4s 62us/sample - loss: 1.1145 - acc: 0.7684 - val\_loss: 1.0179  
- val\_acc: 0.7978

Epoch 5/20

60000/60000 [=====] - 4s 65us/sample - loss: 0.9464 - acc: 0.7982 - val\_loss: 0.8974  
- val\_acc: 0.8056

Epoch 6/20

60000/60000 [=====] - 5s 78us/sample - loss: 0.8337 - acc: 0.8151 - val\_loss: 0.7927  
- val\_acc: 0.8197

Epoch 7/20

60000/60000 [=====] - 4s 66us/sample - loss: 0.7437 - acc: 0.8331 - val\_loss: 0.7059  
- val\_acc: 0.8378

Epoch 8/20

60000/60000 [=====] - 5s 78us/sample - loss: 0.6803 - acc: 0.8409 - val\_loss: 0.6451  
- val\_acc: 0.8594

Epoch 9/20

60000/60000 [=====] - 4s 75us/sample - loss: 0.6319 - acc: 0.8479 - val\_loss: 0.6003  
- val\_acc: 0.8564

Epoch 10/20

60000/60000 [=====] - 3s 58us/sample - loss: 0.5896 - acc: 0.8539 - val\_loss: 0.5707  
- val\_acc: 0.8650

Epoch 11/20  
60000/60000 [=====] - 4s 65us/sample - loss: 0.5598 - acc: 0.8615 - val\_loss: 0.5428  
- val\_acc: 0.8699

Epoch 12/20  
60000/60000 [=====] - 4s 61us/sample - loss: 0.5270 - acc: 0.8666 - val\_loss: 0.5002  
- val\_acc: 0.8772

Epoch 13/20  
60000/60000 [=====] - 4s 68us/sample - loss: 0.5020 - acc: 0.8711 - val\_loss: 0.4835  
- val\_acc: 0.8810

Epoch 14/20  
60000/60000 [=====] - 3s 53us/sample - loss: 0.4807 - acc: 0.8740 - val\_loss: 0.4748  
- val\_acc: 0.8764

Epoch 15/20  
60000/60000 [=====] - 3s 56us/sample - loss: 0.4715 - acc: 0.8761 - val\_loss: 0.4647  
- val\_acc: 0.8828

Epoch 16/20  
60000/60000 [=====] - 4s 61us/sample - loss: 0.4504 - acc: 0.8814 - val\_loss: 0.4375  
- val\_acc: 0.8869

Epoch 17/20  
60000/60000 [=====] - 3s 57us/sample - loss: 0.4378 - acc: 0.8854 - val\_loss: 0.4430  
- val\_acc: 0.8872


Epoch 18/20  
60000/60000 [=====] - 4s 66us/sample - loss: 0.4322 - acc: 0.8865 - val\_loss: 0.4360  
- val\_acc: 0.8881

Epoch 19/20  
60000/60000 [=====] - 4s 66us/sample - loss: 0.4200 - acc: 0.8884 - val\_loss: 0.4168  
- val\_acc: 0.8837

Epoch 20/20  
60000/60000 [=====] - 4s 70us/sample - loss: 0.4066 - acc: 0.8897 - val\_loss: 0.4167  
- val\_acc: 0.8916

```
In [15]: # Plot training & validation accuracy values
def plot_keras_history(h):
    plt.plot(h.history['acc'])
    plt.plot(h.history['val_acc'])
    plt.title('Model accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Test'], loc='upper left')
    plt.show()

plot_keras_history(history)
```



## Let's examine few good and bad cases for which our model works and doesn't work

```
In [16]: y_pred = model_1.predict_classes(x_test.reshape((10000,28*28)))  ## predict_classes can directly give us predicted classes
print(y_pred.shape)

(10000,)
```

```
In [18]: ## First 20 examples
n = 20
i = 0
good_cases = []
for y_hat , y in zip(y_pred, y_test):
    if y_hat == y:
        good_cases.append(i)
        if len(good_cases) == n:
            break
    i += 1

i = 0
bad_cases = []
for y_hat , y in zip(y_pred, y_test):
    if y_hat != y:
        bad_cases.append(i)
        if len(bad_cases) == n:
            break
    i += 1

print(good_cases , "\n", bad_cases)

[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 19, 20]
[18, 33, 62, 77, 87, 92, 104, 124, 125, 149, 150, 151, 172, 185, 187, 195, 217, 233, 235, 241]
```

Now Let's visualize them

```
In [35]: def show_mnist(data, y_pred, y_true, n_rows=0, n_cols=0, title=""):
n = len(data)
assert n==n_rows*n_cols , "n_rows*n_cols should be equal to nun of samples"

fig, axs = plt.subplots(n_rows, n_cols, sharex=True, sharey=True, squeeze=False, figsize=(20,15))


fig.suptitle(title)





for i in range(n_rows):
    for j in range(n_cols):
        m = i*n_cols + j
        img = data[m]
        sub_title = "y_true:{}, predicted:{}".format(y_true[m], y_pred[m])
        axs[i,j].imshow(img, cmap="gray")
        axs[i,j].title.set_text(sub_title)
```



```
In [36]: show_mnist(x_test[good_cases],  
                  y_pred[good_cases],  
                  y_test[good_cases],  
                  n_rows=5, n_cols=4,  
                  title= "good_cases")
```


good\_cases









```
In [37]: show_mnist(x_test[bad_cases],
                  y_pred[bad_cases],
                  y_test[bad_cases],
                  n_rows=5, n_cols=4,
                  title= "bad_cases")
```

bad\_cases





## Second Try

Let's do a second try to improve the performance of this model

**Let's examine the data once again**

```
In [39]: np.max(x_train) , np.max(x_test)
```


```
Out[39]: (255, 255)
```

- As shown above the image contains gray scale values in integer format
- let's scale them back to -1 to +1 range and try again

```
In [46]: x_train_new = x_train / 255  
x_test_new = x_test / 255
```

```
In [47]: plt.imshow(x_train_new[0] , cmap="gray")
```

```
Out[47]: <matplotlib.image.AxesImage at 0x20137e13278>
```



```
In [48]: ## Use a special layer called dropout to avoid over fitting  
from tensorflow.keras.layers import Dropout
```

```
In [49]: def create_model_2():
model = Sequential() # MLP Model
model.add(Dense(64, activation="relu", input_shape=(28*28,)))
model.add(Dropout(0.25))
model.add(Dense(64, activation="relu"))
model.add(Dropout(0.25))
model.add(Dense(10, activation="softmax"))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=["accuracy"])
model.summary()
return model

model_2 = create_model_2()
```

Layer (type)	Output Shape	Param #
=====		
dense_10 (Dense)	(None, 64)	50240
dropout_2 (Dropout)	(None, 64)	0
dense_11 (Dense)	(None, 64)	4160
dropout_3 (Dropout)	(None, 64)	0
dense_12 (Dense)	(None, 10)	650
=====		
Total params: 55,050		
Trainable params: 55,050		
Non-trainable params: 0		
=====		

**Let's train Model 2**

In [50]: *# Make sure to use the new rescaled datasets x\_train\_new and x\_test\_new*

```
history_2 = model_2.fit(x=x_train_new.reshape((60000, 28*28)),
                        y=y_train_vecs,
                        batch_size=64,
                        epochs=20,
                        verbose=1,
                        validation_data=(x_test_new.reshape((10000, 28*28)), y_test_vecs))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 118us/sample - loss: 0.4971 - acc: 0.8497 - val\_loss: 0.1815  
- val\_acc: 0.9456

Epoch 2/20

60000/60000 [=====] - 5s 90us/sample - loss: 0.2454 - acc: 0.9272 - val\_loss: 0.1391  
- val\_acc: 0.9579

Epoch 3/20

60000/60000 [=====] - 5s 87us/sample - loss: 0.2024 - acc: 0.9389 - val\_loss: 0.1274  
- val\_acc: 0.9603

Epoch 4/20

60000/60000 [=====] - 5s 91us/sample - loss: 0.1808 - acc: 0.9462 - val\_loss: 0.1241  
- val\_acc: 0.9641

Epoch 5/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.1629 - acc: 0.9520 - val\_loss: 0.1098  
- val\_acc: 0.9680

Epoch 6/20

60000/60000 [=====] - 6s 96us/sample - loss: 0.1507 - acc: 0.9542 - val\_loss: 0.1064  
- val\_acc: 0.9682

Epoch 7/20

60000/60000 [=====] - 6s 92us/sample - loss: 0.1411 - acc: 0.9566 - val\_loss: 0.1006  
- val\_acc: 0.9698

Epoch 8/20

60000/60000 [=====] - 5s 83us/sample - loss: 0.1355 - acc: 0.9591 - val\_loss: 0.0942  
- val\_acc: 0.9729

Epoch 9/20

60000/60000 [=====] - 5s 88us/sample - loss: 0.1268 - acc: 0.9615 - val\_loss: 0.0932  
- val\_acc: 0.9725

Epoch 10/20

60000/60000 [=====] - 5s 85us/sample - loss: 0.1223 - acc: 0.9623 - val\_loss: 0.0958  
- val\_acc: 0.9719

Epoch 11/20


60000/60000 [=====] - 5s 82us/sample - loss: 0.1150 - acc: 0.9645 - val\_loss: 0.0890  
- val\_acc: 0.9741

Epoch 12/20

60000/60000 [=====] - 5s 87us/sample - loss: 0.1142 - acc: 0.9643 - val\_loss: 0.0852  
- val\_acc: 0.9755  
Epoch 13/20  
60000/60000 [=====] - 5s 84us/sample - loss: 0.1099 - acc: 0.9664 - val\_loss: 0.0902  
- val\_acc: 0.9733  
Epoch 14/20  
60000/60000 [=====] - 5s 87us/sample - loss: 0.1074 - acc: 0.9658 - val\_loss: 0.0863  
- val\_acc: 0.9746  
Epoch 15/20  
60000/60000 [=====] - 5s 86us/sample - loss: 0.1043 - acc: 0.9681 - val\_loss: 0.0928  
- val\_acc: 0.9735  
Epoch 16/20  
60000/60000 [=====] - 5s 89us/sample - loss: 0.1019 - acc: 0.9678 - val\_loss: 0.0868  
- val\_acc: 0.9742  
Epoch 17/20  
60000/60000 [=====] - 5s 89us/sample - loss: 0.0996 - acc: 0.9693 - val\_loss: 0.0880  
- val\_acc: 0.9729  
Epoch 18/20  
60000/60000 [=====] - 6s 107us/sample - loss: 0.0952 - acc: 0.9699 - val\_loss: 0.0905  
- val\_acc: 0.9734  
Epoch 19/20  
60000/60000 [=====] - 5s 91us/sample - loss: 0.0911 - acc: 0.9722 - val\_loss: 0.0853  
- val\_acc: 0.9750  
Epoch 20/20  
60000/60000 [=====] - 5s 87us/sample - loss: 0.0936 - acc: 0.9699 - val\_loss: 0.0876  
- val\_acc: 0.9747



```
In [52]: plot_keras_history(history_2)
```



## WOW ! 97.5% accuracy in validation

- Now with these changes we have got 97.5% accuracy . **But still we missed 2.5%** . We would examine them latter
- Model\_2 inorporated following changes
  - Activation function of layes chnaged to "relu"
  - Layer dims increased to 64

- Dropout added to avoid overfitting
- Data rescaled

## Let's examine bad cases for which even our updated model doesn't work

```
In [54]: y_pred_2 = model_2.predict_classes(x_test_new.reshape((10000,28*28)))  
print("y_pred_2 : ", y_pred_2.shape)
```

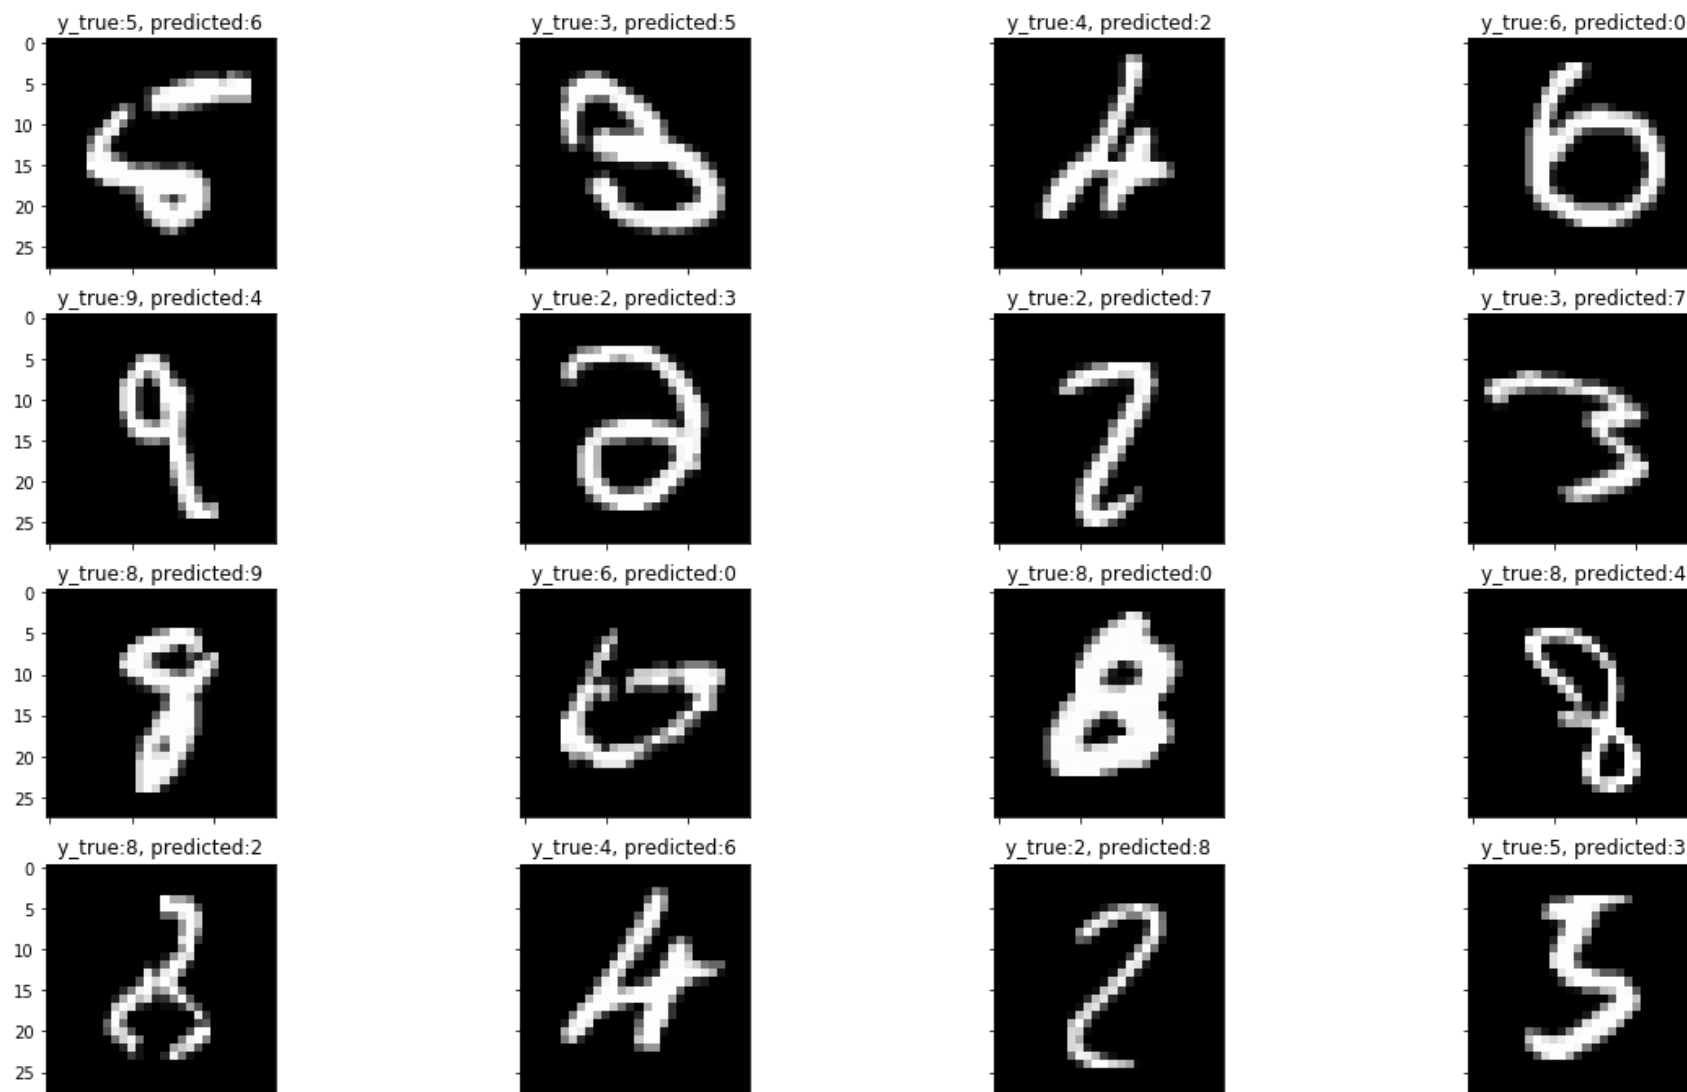
```
y_pred_2 : (10000,)
```




```
In [56]: ## First 20 bad examples  
n = 20  
i = 0  
bad_cases_2 = []  
for y_hat , y in zip(y_pred_2, y_test):  
    if y_hat != y:  
        bad_cases_2.append(i)  
        if len(bad_cases_2) == n:  
            break  
    i += 1  
  
print(bad_cases_2)
```

```
[8, 18, 247, 259, 264, 318, 321, 381, 435, 445, 495, 543, 582, 610, 613, 674, 691, 717, 720, 740]
```

```
In [58]: show_mnist(x_test[bad_cases_2],
                    y_pred_2[bad_cases_2],
                    y_test[bad_cases_2],
                    n_rows=5, n_cols=4,
                    title= "bad_cases_2")
```

bad\_cases\_2





**Those are really tough .. but we can achieve > 99% accuracy with CNNs !**

Done !